

Dell EMC VxFlex OS - CSI Driver for Kubernetes

Version 1.0.0.0

Product Guide

305-XXX-XXX

01

Copyright © 2019 Dell Inc. or its subsidiaries. All rights reserved.

Published April 2019

Dell believes the information in this publication is accurate as of its publication date. The information is subject to change without notice.

THE INFORMATION IN THIS PUBLICATION IS PROVIDED “AS-IS.” DELL MAKES NO REPRESENTATIONS OR WARRANTIES OF ANY KIND WITH RESPECT TO THE INFORMATION IN THIS PUBLICATION, AND SPECIFICALLY DISCLAIMS IMPLIED WARRANTIES OF MERCHANTABILITY OR FITNESS FOR A PARTICULAR PURPOSE. USE, COPYING, AND DISTRIBUTION OF ANY DELL SOFTWARE DESCRIBED IN THIS PUBLICATION REQUIRES AN APPLICABLE SOFTWARE LICENSE.

Dell, EMC, and other trademarks are trademarks of Dell Inc. or its subsidiaries. Other trademarks may be the property of their respective owners.
Published in the USA.

Dell EMC
Hopkinton, Massachusetts 01748-9103
1-508-435-1000 In North America 1-866-464-7381
www.DellEMC.com

CONTENTS

Chapter 1	Introduction	5
	Product overview.....	6
Chapter 2	Installing 1.0 driver	7
	Installation overview.....	8
	Prerequisites.....	8
	Enable Kubernetes feature gates.....	8
	Configure Docker service.....	10
	Install Helm package manager.....	10
	Install VxFlex OS Storage Data Client.....	11
	Install the VxFlex OS - CSI driver.....	11
Chapter 3	Testing driver	15
	Test deploying a simple pod with VxFlex OS storage.....	16
	Test creating snapshots.....	17
	Test restoring from a snapshot.....	18
Chapter 4	Troubleshooting	19
	Troubleshooting.....	20

CONTENTS

CHAPTER 1

Introduction

- [Product overview](#)6

Product overview

The VxFlex OS - CSI driver is a plug-in that is installed into Kubernetes to provide persistent storage using Dell EMC VxFlex OS storage system.

The VxFlex OS - CSI driver and Kubernetes communicate using the Container Storage Interface protocol. VxFlex OS - Container Storage Interface (CSI) 1.0 Driver for Kubernetes 1.0 driver conforms to CSI specification v1.0 and compatible with Kubernetes versions 1.13.1, 1.13.2 and 1.13.3.

Features of VxFlex OS - CSI driver

The VxFlex OS - CSI driver supports the following features:

- Persistent volume (PV) capabilities - create, list, delete, and create-from-snapshot
- Dynamic and static PV provisioning
- Snapshot capabilities - create, delete, and list
- Volume mount as *ext4* or *xf*s file system on the worker node
- Supports the following access modes:
 - multi-node-single-writer
 - single-node-writer
 - single-node-reader-only
 - multi-node-reader-only
 - multi-node-single-writer
- Volume prefix for easy LUN identification
- Supports HELM charts installer
- Supports CSI 1.0
- Supports Kubernetes version 1.13.1, 1.13.2 and 1.13.3
- Compatible with VxFlex OS version 3.0 with medium granularity storage pools.

The VxFlex OS - CSI driver depends on the following libraries:

- *gofsutil*
- *gocsi*
- *goscaleio*
- *godog*
- *protobuf*
- *logrus*
- *gosync*
- *grpc*

These libraries are used in the current version of VxFlex OS - CSI driver.

CHAPTER 2

Installing 1.0 driver

This chapter contains the following sections:

- [Installation overview](#) 8
- [Prerequisites](#) 8
- [Install the VxFlex OS - CSI driver](#) 11

Installation overview

The Helm chart installs the VxFlex OS - CSI driver using a shell script. This script installs the CSI driver container image along with the required Kubernetes sidecar containers.

The controller section of the Helm chart installs the following components in a *Stateful Set* in the namespace *vxflexos*:

- VxFlex OS - CSI driver
- Kubernetes Provisioner, which provisions the provisioning volumes
- Kubernetes Attacher, which attaches the volumes to the containers
- Kubernetes Snapshotter, which provides snapshot support

The node section of the Helm chart installs the following component in a *Daemon Set* in the namespace *vxflexos*:

- VxFlex OS - CSI driver
- Kubernetes Registrar, which handles the driver registration

Prerequisites

Before you install VxFlex OS - CSI driver, verify the requirements that are mentioned in this topic are installed and configured.

- Install Kubernetes. The VxFlex OS - CSI driver works with Kubernetes version 1.13.1 or later.
- [Enable the Kubernetes feature gates](#)
- [Configure Docker service](#)
- [Install Helm and Tiller with a service account](#)
- [Install VxFlex OS SDC](#)

Enable Kubernetes feature gates

The Kubernetes feature gates must be enabled before installing the VxFlex OS - CSI driver.

The [Feature Gates section](#) of Kubernetes home page lists the Kubernetes feature gates. The following Kubernetes feature gates must be enabled:

- VolumeSnapshotDataSource
- KubeletPluginsWatcher
- CSINodeInfo
- CSIDriverRegistry
- BlockVolume
- CSIBlockVolume

Procedure

1. On each master and node of Kubernetes, edit `/var/lib/kubelet/config.yaml` and add the following lines at the end to set feature-gate settings for the kubelets:

```
/var/lib/kubelet/config.yaml
VolumeSnapshotDataSource: true
```



```
KubeletPluginsWatcher: true
CSINodeInfo: true
CSIDriverRegistry: true
BlockVolume: true
CSIBlockVolume: true
```

2. On the master, set the feature gate settings of the *kube-apiserver.yaml* file as follows:

```
/etc/kubernetes/manifests/kube-apiserver.yaml - --feature-
gates=VolumeSnapshotDataSource=true,KubeletPluginsWatcher=true
,CSINodeInfo=true,CSIDriverRegistry=true,BlockVolume=true,CSIB
lockVolume=true
```

3. On the master, set the feature gate settings of the *kube-controller-manager.yaml* file as follows:

```
/etc/kubernetes/manifests/kube-controller-manager.yaml - --
feature-
gates=VolumeSnapshotDataSource=true,KubeletPluginsWatcher=true
,CSINodeInfo=true,CSIDriverRegistry=true,BlockVolume=true,CSIB
lockVolume=true
```

4. On the master, set the feature gate settings of the *kube-scheduler.yaml* file as follows:

```
/etc/kubernetes/manifests/kube-scheduler.yaml - --feature-
gates=VolumeSnapshotDataSource=true,KubeletPluginsWatcher=true
,CSINodeInfo=true,CSIDriverRegistry=true,BlockVolume=true,CSIB
lockVolume=true
```

5. On each node, edit the variable **KUBELET_KUBECONFIG_ARGS** of */etc/systemd/system/kubelet.service.d/kubeadm.conf* file as follows:

```
Environment="KUBELET_KUBECONFIG_ARGS=--bootstrap-
kubeconfig=/etc/kubernetes/bootstrap-kubelet.conf --
kubeconfig=/etc/kubernetes/kubelet.conf --allow-
privileged=true --
featuregates=VolumeSnapshotDataSource=true,KubeletPluginsWatch
er=true,CSINodeInfo=true,CSIDriverRegistry=true,BlockVolume=tr
ue,CSIBlockVolume=true"
```

6. Restart the kublet with **systemctl daemon-reload** and **systemctl restart kubelet** on all nodes.

Configure Docker service

The mount propagation in Docker must be configured on all Kubernetes nodes before installing the VxFlex OS - CSI driver.

Procedure

1. Edit the service section of `/etc/systemd/system/multi-user.target.wants/docker.service` file as follows:

```
docker.service
[Service]
...
MountFlags=shared
```

2. Restart the docker service with `systemctl daemon-reload` and `systemctl restart docker` on all the nodes.

Install Helm package manager

Install the Helm and Tiller package manager on the master node.

To install Helm, run the following commands:

Procedure

1. Run `curl https://raw.githubusercontent.com/helm/helm/master/scripts/get > get_helm.sh`
2. Run `chmod 700 get_helm.sh`
3. Run `./get_helm.sh`
4. Run `helm.init`
5. Run `helm version` to test the Helm installation.
6. Set up a service account for Tiller as follows:
 - a. Add the following snippet to the `rbac-config.yaml` file:

```
apiVersion: v1
kind: ServiceAccount
metadata:
  name: tiller
  namespace: kube-system
---
kind: ClusterRoleBinding
apiVersion: rbac.authorization.k8s.io/v1beta1
metadata:
  name: tiller-clusterrolebinding
subjects:
- kind: ServiceAccount
  name: tiller
  namespace: kube-system
roleRef:
  kind: ClusterRole
  name: cluster-admin
  apiGroup: ""
```

- b. Run `kubectl create -f rbac-config.yaml` to create the service account.

7. Apply the service account to Tiller using `helm init --upgrade --service-account tiller`

Install VxFlex OS Storage Data Client

Install the VxFlex OS Storage Data Client (SDC) on all Kubernetes nodes.

For detailed VxFlex OS installation procedure, see the *Dell EMC VxFlex OS Deployment Guide*. Install the VxFlex OS SDC as follows:

Procedure

1. Download the VxFlex OS SDS from [Dell EMC Online support](#). The filename is EMC-ScaleIO-sdc-*.rpm, where * is the SDC name corresponding to the VxFlex OS installation version.
2. Export the shell variable `MDM_IP` in a comma-separated list. This list contains the IP addresses of the MDMs.

```
export MDM_IP=10.254.16.58,10.254.16.62
```
3. Install the SDC with the command `rpm -iv ./EMC-ScaleIO-sdc-*.x86_64.rpm`, where * is the SDC name corresponding to the VxFlex OS installation version.

Install the VxFlex OS - CSI driver

Install the VxFlex OS - CSI driver using this procedure.

Before you begin

- You must have the downloaded files, including the Helm chart from the source repository at github.com/dell/csi-vxflexos, ready for this procedure.
- In the top level helm directory, there should be two shell scripts, `install.vxflexos` and `uninstall.vxflexos`. These scripts handle some of the pre and post operations that cannot be performed in the helm chart, such as creating Custom Resource Definitions (CRDs), if needed.
- You must create a Kubernetes secret with your VxFlex OS username and password. Use the `secret.yaml` file to create the secret with the following values to match the default installation parameters:
 - Name: `vxflexos-creds`
 - Namespace: `vxflexos`

Note

- Create the namespace using `kubectl create namespace vxflexos`.
 - For more information about creating a Kubernetes secret, see: [Kubernetes documentation: Overview of Secrets](#)
-

Procedure

1. Collect information from the VxFlex OS SDC (Storage Data Client) by executing the `get_vxflexos_info.sh` script located in the top-level helm directory.

This script shows the *VxFlex OS system ID* and *MDM IP* addresses. Make a note of the value for these parameters as they must be entered in the `myvalues.yaml` file.

Note

Your SDC might have multiple VxFlex OS systems registered. Ensure that you choose the correct values.

2. Copy the `csi-vxflexos/values.yaml` into a file in the same directory as the `install.yaml` named `myvalues.yaml`, to customize settings for installation.
 3. Edit `myvalues.yaml` to set the following parameters for your installation:
 - Set the `systemName` string variable to the VxFlex OS system name or system ID. This value was obtained by running the `get_vxflexos_info.sh` script in Step 1 of this procedure.
 - Set the `restGateway` string variable to the URL of your system's REST API Gateway. You can obtain this value from the VxFlex OS administrator.
 - Set the `storagePool` string variable to a default (already existing) storage pool name in your VxFlex OS system.
-

Note

New storage pools can be created in VxFlex OS UI and CLI utilities.

- Set the `mdmIP` string variable to a comma separated list of MDM IP addresses.
- Set the `volumeNamePrefix` string variable so that volumes created by the driver have a default prefix. If one VxFlex OS system is servicing several different Kubernetes installations or users, these prefixes help you distinguish them.
- The `controllerCount` variable is used by advanced users to deploy multiple controller instances. The specified default value 1 is designed to work as expected. You can modify the value of this variable to set the desired number of CSI controller replicas.
- Set the `enablelistvolumesnapshot` variable `false` unless instructed otherwise, by Dell EMC support. It causes snapshots to be included in the CSI operation ListVolumes.
- The Helm charts create a Kubernetes `StorageClass` while deploying the VxFlex OS - CSI Driver. The `StorageClass` section includes following variables:
 - The `name` string variable defines the name of the Kubernetes storage class that the Helm charts will create. For example, the "vxflexos" base name will be used to generate names such as "vxflexos" and "vxflexos-xfs".
 - The `reclaimPolicy` string variable defines whether the volumes will be retained or deleted when the assigned pod is destroyed. The valid values for this variable are `Retain` or `Delete`.
 - The `isDefault` variable (valid values for this variable are `true` or `false`) will set the newly created storage class as default for Kubernetes.

Note

- Set this value to “true” only if you expect VxFlex OS to be your principle storage provider, as it will be used in *PersistentVolumeClaims* where no *storageclass* is provided. After installation, you can add custom storage classes if desired.
 - All strings must be contained within double quotes.
-

4. Run the `sh install.vxflexos` command to proceed with the installation.

A successful installation should emit messages that look similar to the following samples:

```
[root@k8s113a-10-247-102-211 helm]# sh install.vxflexos
NAME:      vxflexos
LAST DEPLOYED: Thu Jan 24 18:28:35 2019
NAMESPACE: vxflexos
STATUS: DEPLOYED

... omitted lots here ...

NAME                READY    STATUS    RESTARTS   AGE
vxflexos-controller-0  4/4      Running    0           20s
vxflexos-node-r5kdt    2/2      Running    0           20s
vxflexos-node-tq5tj    2/2      Running    0           20s
CSIDrivers:
NAME                AGE
vxflexos            21s
CSINodeInfos:
NAME                                AGE
k8s113a-10-247-102-213.lss.emc.com  21h
k8s113a-10-247-102-215.lss.emc.com  21h
StorageClasses:
NAME                PROVISIONER    AGE
vxflexos (default)  csi-vxflexos   21s
vxflexos-xfs        csi-vxflexos   21s
VolumeSnapshotClasses:
NAME                AGE
vxflexos-snapclass  2d
```

Results

At the end of the script, the `kubectl get pods -n vxflexos` is called to GET the status of the pods and you will see the following:

- Vxflex-controller-0 with 4/4 containers ready, and status displayed as `Running`.
- Agent pods with 2/2 containers and the status displayed as `Running`.

Finally, the script lists the created *storageclasses* such as, "vxflexos" and "vxflexos-xfs". Additional storage classes can be created for different combinations of file system types and VxFlex OS storage pools. The script also creates *volumesnapshotclasses* such as, "vxflexos-snapclass" and other snapshots classes.

CHAPTER 3

Testing driver

This chapter contains the following sections:

- [Test deploying a simple pod with VxFlex OS storage](#)..... 16
- [Test creating snapshots](#)..... 17
- [Test restoring from a snapshot](#)..... 18

Test deploying a simple pod with VxFlex OS storage

Test the deployment workflow of a simple pod on VxFlex OS storage.

Before you begin

In the source code, there is a directory that contains examples of how you can use the driver. To use these examples, you must create a "test" namespace, using "kubectl create namespace test", before you can start testing.

The *starttest.sh* script is located in the `csi-vxflexos/test/helm` directory. This script is used in the following procedure to deploy helm charts that test the deployment of a simple pod.

Procedure

1. Navigate to the `test/helm` directory, which contains the *starttest.sh* and the *2vols* directories.

This directory contains a simple Helm chart that will deploy a pod that uses two VxFlex OS volumes.

2. Run the `sh starttest.sh 2vols` command to deploy the pod.

You should see the following:

```
Normal      Pulled          38s          kubelet,
k8s113a-10-247-102-215.1ss.emc.com Successfully pulled image
"docker.io/centos:latest"
Normal      Created          38s
kubelet, k8s113a-10-247-102-215.1ss.emc.com Created container
Normal      Started          38s
kubelet, k8s113a-10-247-102-215.1ss.emc.com Started container
/dev/scinib      8125880    36852    7653216    1% /data0
/dev/scinia      16766976   32944    16734032   1% /data1
/dev/scinib on /data0 type ext4 (rw,relatime,data=ordered)
/dev/scinia on /data1 type xfs
(rw,relatime,attr2,inode64,noquota)
```

3. To stop the test, run `sh stoptest.sh`.

This script will delete the pods and the volumes depending on the retention setting you have configured.

Results

An outline of this workflow is described below:

1. The *2vols* helm chart contains two PersistentVolumeClaim definitions, one in *pvc0.yaml*, and the other in *pvc1.yaml*. They are referenced by the *test.yaml* which creates the pod. The contents of the *Pvc0.yaml* file are described below:

```
kind: PersistentVolumeClaim
apiVersion: v1
metadata:
  name: pvc0
  namespace: test
spec:
  accessModes:
```



```
- ReadWriteOnce
volumeMode: Filesystem
resources:
  requests:
    storage: 8Gi
storageClassName: vxflexos
```

2. The *volumeMode: Filesystem* requires a mounted file system and the *resources.requests.storage* of 8Gi requires an 8 GB file. In this case, the *storageClassName: vxflexos* directs the system to use one of the pre-defined storage classes created by the VxFlex OS - CSI installation process. This step yields a mounted *ext4* file system. You can see the storage class definitions in the VxFlex OS installation helm chart files *storageclass.yaml* and *storageclass-xfs.yaml*.
3. If you compare *pvol0.yaml* and *pvol1.yaml*, you will find that the latter uses a different storage class; *vxflexos-xfs*. This class gives you an *xfs* file system.
4. To see the volumes you created, run `kubectl get persistentvolumeclaim -n test` and `kubectl describe persistentvolumeclaim -n test`.

Note

For more information about Kubernetes objects like *StatefulSet* and *PersistentVolumeClaim* see [Kubernetes documentation: Concepts](#).

Test creating snapshots

Test the workflow for snapshot creation.

Procedure

1. Start the *2vols* container and leave it running.
2. Run the `snaptest.sh` shell script.

This will create a snapshot of each of the volumes in the container using *VolumeSnapshot* objects defined in *snap1.yaml* and *snap2.yaml*. The following are the contents of *snap1.yaml*:

```
apiVersion: snapshot.storage.k8s.io/v1alpha1
kind: VolumeSnapshot
metadata:
  name: pvol0-snap1
  namespace: test
spec:
  snapshotClassName: vxflexos-snapclass
  source:
    name: pvol0
    kind: PersistentVolumeClaim
```

Results

The *snaptest.sh* script will create a snapshot using the definitions in the *snap1.yaml* file. The *spec.source* section contains the volume that will be snapped. For example, if the volume to be snapped is *pvol0*, then the created snapshot is named *pvol0-snap1*.

Notice that this *VolumeSnapshot* class has a reference to a *snapshotClassName: vxflexos-snapclass*. The VxFlex OS - CSI driver installation creates this class as its default snapshot class. You can see its definition in the installation directory file *volumesnapshotclass.yaml*.

Note

The *snaptest.sh* shell script creates the snapshots, describes them, and then deletes them. You can see your snapshots using *kubectl get volumesnapshot -n test*.

Test restoring from a snapshot

Test the restore operation workflow to restore from a snapshot.

Before you begin

Ensure that you have stopped any previous test instance before performing this procedure.

To test the restore operation from a snapshot:

Procedure

1. Run the `snaprestoretest.sh` shell script.

This script deploys the *2vols* example, creates a snap of *pvol0*, and then updates the deployed helm chart from the updated directory *2vols+restore*. This then adds an additional volume that is created from the snapshot.

Results

An outline of this workflow is described below:

1. The snapshot is taken using *snap1.yaml*.
2. *Helm* is called to upgrade the deployment with a new definition, which is found in the *2vols+restore* directory. The `helm/templates` directory contains the newly created *createFromSnap.yaml* file. The script then creates a *PersistentVolumeClaim*, which is a volume that is dynamically created from the snapshot. Then the helm deployment is upgraded to contain the newly created third volume. In other words, when the *snaprestoretest.sh* creates a new volume with data from the snapshot, the restore operation is tested. The contents of the *createFromSnap.yaml* are described below:

```
apiVersion: v1
kind: PersistentVolumeClaim
metadata:
  name: restorepvc
  namespace: test
spec:
  storageClassName: vxflexos
  dataSource:
    name: pvol0-snap1
    kind: VolumeSnapshot
    apiGroup: snapshot.storage.k8s.io
  accessModes:
    - ReadWriteOnce
  resources:
    requests:
      storage: 8Gi
```

Note

The *spec.dataSource* clause, specifies a source "VolumeSnapshot" named *pvol0-snap1* which matches the snapshot's name in *snap1.yaml*.

CHAPTER 4

Troubleshooting

This chapter contains the following topics:

- [Troubleshooting](#) 20

Troubleshooting

The following table lists the VxFlex OS CSI driver installation troubleshooting scenarios:

Symptoms	Prevention, resolution, or workaround
<p>The installation fails with the following error message:</p> <pre>Node xxx does not have the SDC installed</pre>	<p>Install the SDC on listed nodes. The SDC must be installed on all Kubernetes master and minion nodes.</p>
<p>When you run the command <code>kubectl describe pods vxflexos-controller-0 -n vxflexos</code>, the system indicates that the driver image could not be loaded.</p>	<p>Edit the <i>daemon.json</i> file found in the registry location and add</p> <pre>{ "insecure-registries" : ["hostname.cloudapp.net:5000"] }</pre> <p>.</p>
<p>The <code>kubectl logs vxflexos-controller-0 -n vxflexos</code> logs shows that the driver is not authenticated.</p>	<p>Check the username, password, and the gateway IP address for the VxFlex OS system.</p>
<p>The <code>kubectl logs vxflexos-controller-0 -n vxflexos</code> logs shows that the system ID is incorrect.</p>	<p>Use the <code>get_vxflexos_info.sh</code> to find the correct system ID. Add the system ID to <i>myvalues.yaml</i> script.</p>